

1 Newton's method for computing least squares

(a) Hessian (H) is a matrix of second derivatives. We can start from the expression for a derivative of cost function $J(\Theta)$ as was presented in the lecture notes:

$$\frac{\partial}{\partial \Theta_j} J(\Theta) = (h_{\Theta}(x) - y)x_j \rightarrow \nabla_{\Theta} J(\Theta) = X^T X \Theta - X^T \vec{y}$$

where $h_{\Theta}(x) = \sum_{i=0}^n \Theta_i x_i$, and differentiate it second time wrt Θ_k :

$$\frac{\partial^2}{\partial \Theta_j \partial \Theta_k} J(\Theta) = x_j x_k$$

This can be also expressed in a matrix form: $H = X^T X$.

(b) The step in the Newton's method is defined:

$$\Theta := \Theta - \frac{f(\Theta)}{f'(\Theta)}$$

when one looks for roots of $f(\Theta)$. Since we want to minimize $J(\Theta)$, we want $\nabla_{\Theta} J(\Theta) = 0$, and thus we use the modified version expressed on matrices and vectors:

$$\Theta := \Theta - H^{-1} \nabla_{\Theta} J(\Theta)$$

. Let us now look how looks Θ after first step (Θ_1) taking some Θ_0 as a starting point, and implementing the expressions found in the previous exercise:

$$\Theta_1 = \Theta_0 - (X^T X)^{-1} (X^T X \Theta_0 - X^T \vec{y}) = \Theta_0 - \Theta_0 + (X^T X)^{-1} X^T \vec{y} = (X^T X)^{-1} X^T \vec{y}$$

As we can see our expression do not depend on the choice of Θ_0 , and we arrive to the solution in one step.

2 Locally-weighted logistic regression

(a) The program written for this exercise is available in `q2/` directory, here we present the function `lwlr.m`:

```
function y = lwlr(X_train, y_train, x, tau)

%% Locally weighted logistic regression
% the value of parameter lambda given in exercise
lam = 0.0001;
%% m the sample size
m = length(y_train);
```

```

%% Update of Theta in Newton's method Th := Th + l'(Th)/l''(Th)
%% Initialise Th as a zero vector:
Th = [0;0];
Grad = [1;1];

%% to avoid too many iterations
iter = 0;
%% building vector of weights for the training set
weights = exp(-sum((X_train - repmat(x', m, 1)).^2, 2) / (2 * tau^2));

while (norm(Grad) > 0.0001 && iter < 10000)
    iter = iter+1;
%% vector of h_Theta from X
    h = 1.0 ./ (1.0 + exp(-sum(X_train .* Th', 2)));

%% auxiliary variables z and D for gradient and hessian calculation
    z = weights .* (y_train - h);
    D = diag(- weights .* (1.0 - h) .* h);

    %% Hessian expression H = XT DX I
    Hess = X_train' * D * X_train - lam * eye(2);
    %% Grad expression l() = XT z
    Grad = X_train' * z - lam * Th;
    %% Update of Theta in Newton's method Th := Th + l'(Th)/l''(Th)
    Th = Th - Hess \ Grad;
endwhile
%% make prediction for y
if (1.0 / (1.0 + exp(-x' * Th)) > 0.5)
    y = 1;
else
    y = 0;
endif

```

(b) There will be pics at some point...

3 Multivariate least squares

(a) We have a cost function:

$$J(\Theta) = \frac{1}{2} \sum_{i=1}^m \sum_{j=1}^p ((\Theta^T x^{(i)})_j - y_j^{(i)})^2$$

. To express this in a vector/matrix form let us consider building blocks: the design matrix $X \in \mathbb{R}^{m \times n}$, the target matrix $Y \in \mathbb{R}^{m \times p}$ and the parameters matrix $\Theta \in \mathbb{R}^{n \times p}$. The cost function

expressed with these matrices:

$$J(\Theta) = \frac{1}{2} \text{tr}(X\Theta - Y)^T(X\Theta - Y)$$

where matrices have right dimensions that the operations are well defined and the trace is applied to take the sum of diagonal elements (\sum_j in first expression).

(b) The minimization of $J(\Theta)$ requires $\nabla_{\Theta} J(\Theta) = 0$:

$$\nabla_{\Theta} J(\Theta) = \frac{1}{2} \nabla_{\Theta} \text{tr}(X\Theta - Y)^T(X\Theta - Y) = \frac{1}{2} \nabla_{\Theta} \text{tr}(\Theta^T X^T X \Theta - Y^T X \Theta - \Theta^T X^T Y + Y^T Y) =$$

now we use properties of trace (eg. cyclic or $\text{tr} A = \text{tr} A^T$) and of the matrix derivatives as shown in the Sec. 2.1. of lecture notes:

$$= \frac{1}{2} \nabla_{\Theta} \text{tr} \Theta^T X^T X \Theta - \nabla_{\Theta} \text{tr} Y^T X \Theta = X^T X \Theta - X^T Y$$

We identify it with zero and solve for Θ :

$$X^T X \Theta - X^T Y = 0$$

$$X^T X \Theta = X^T Y$$

$$\Theta = (X^T X)^{-1} X^T Y$$

(c) When there is just one value of y , we get the expression (from the lecture notes):

$$\theta = (X^T X)^{-1} X^T \vec{y}$$

. When we consider the p independent target vectors \vec{y}_j where $j = 1, \dots, p$ we can combine them in a matrix $Y = (\vec{y}_1 \dots \vec{y}_p)$ and also the parameters $\theta \rightarrow \theta_j$ will arrive into the final parameter matrix $\Theta = [\theta_1 \dots \theta_p]$.

4 Naive Bayes

(a)

5 Exponential family and the geometric distribution

(a) Geometric distribution given by:

$$p(y; \phi) = (1 - \phi)^{(y-1)} \phi \quad \text{for } y = 1, 2, 3, \dots$$

and the Exponential Family form:

$$p(y; \eta) = b(y) \exp(\eta^T T(y) - a(\eta))$$

where $T(y)$ is a sufficient statistic, usually $T(y) = y$, and η is natural parameter. Let us play with the form of the geometric distribution to allow the identification of the Exponential Family form:

$$p(y; \phi) = \exp(\log((1 - \phi)^{y-1} \phi)) = \exp((y - 1) \log(1 - \phi) + \log(\phi)) = \exp(y \log(1 - \phi) - \log(\frac{\phi}{1 - \phi}))$$

We can now identify:

$$\begin{aligned} T(y) &= y \\ b(y) &= 1 \\ \eta^T &= \log(1 - \phi) \end{aligned}$$

from this we can get: $\phi = 1 - \exp(\eta)$,

$$a(\eta) = -\log(\frac{\phi}{1 - \phi}) = -\log(\frac{1 - \exp(\eta)}{1 - 1 + \exp(\eta)}) = \log(\frac{1}{1 + \exp(-\eta)})$$

(b) knowing $E[y; \phi] = \frac{1}{\phi}$, we have:

$$g(\eta) = E[T(y); \eta] = E[y; \eta(\phi)] = \frac{1}{phi} = \frac{1}{1 - \exp(\eta)}$$

(c) having $p(y; \phi) = (1 - \phi)^{y-1} \phi$ and $\phi = 1 - \exp(\Theta^T x)$ we get the log-likelihood of a form:

$$\begin{aligned} l(\Theta) &= \log(p(y^{(i)} | x^{(i)}; \Theta)) = \log(1 - \phi)^{y^{(i)} - 1} \phi = \log((\exp(\Theta^T x^{(i)})^{y^{(i)} - 1} (1 - \exp(\Theta^T x))) = (y^{(i)} - 1) \Theta^T x^{(i)} + \log(1 - \exp(\Theta^T x^{(i)})) \\ &= y^{(i)} \Theta^T x^{(i)} + \log(\frac{1 - \exp(\Theta^T x^{(i)})}{\exp(\Theta^T x^{(i)})}) = y^{(i)} \Theta^T x^{(i)} + \log(\exp(-\Theta^T x^{(i)}) - 1) \end{aligned}$$

To derive rule for a stochastic gradient ascent, we need to compute the gradient:

$$\frac{d}{d\Theta_j} l(\Theta) = y^{(i)} x_j^{(i)} + \frac{\exp(-\Theta^T x^{(i)}) (-x_j^{(i)})}{\exp(-\Theta^T x^{(i)}) - 1} = (y^{(i)} - \frac{1}{1 - \exp(\Theta^T x^{(i)})}) x_j^{(i)}$$

and then the rule is:

$$\begin{aligned} \Theta_j &:= \Theta_j + \alpha \frac{\partial l(\Theta)}{\partial \Theta_j} \\ \Theta_j &:= \Theta_j + \alpha (y^{(i)} - \frac{1}{1 - \exp(\Theta^T x^{(i)})}) x_j^{(i)} \end{aligned}$$